

# 431 Class 10

[thomaseLove.github.io/431](https://thomaseLove.github.io/431)

2021-09-23

# Today's Agenda

- 1 Ingesting `dm1000` data using R data set format (`.Rds`)
- 2 Partitioning data into model training/test samples.
- 3 Augmenting a Scatterplot (labeling, size, color) and fitting a simple OLS (linear) model `m1`
- 4 Using `summary()` and `extract_eq()` on a regression model.
- 5 The broom package and `tidy()`, `glance()` and `augment()`
- 6 Calibrating your understanding of R-square a bit
- 7 Assessing Regression Assumptions with Residual Plots
- 8 Making Predictions into the Test Sample
- 9 Assessing Quality of Fit using the Test Sample with mean and maximum absolute prediction error and with RMSPE
- 10 Fitting a Bayesian Linear Model with default priors (`m2`)
- 11 Including Insurance without (`m3`) and with (`m4`) interaction with `dbp` in linear models

# Today's Packages

```
library(broom)
library(equatiomatic) # new today
library(ggrepel) # sort of new today
library(glue) # sort of new today
library(janitor)
library(knitr)
library(magrittr)
library(patchwork)
library(rstanarm) # special today
library(tidyverse)

theme_set(theme_bw())
```

# Data Ingest and Partitioning

# Today's Data

Today, we'll use an R data set (.Rds) to import the dm1000 data.

```
dm1000 <- read_rds("data/dm_1000.Rds")
```

- This allows us to read in the data just as they were last saved in R, including “factoring” and handling of missing data, etc. The function `readRDS()` also works but is a little slower.
- To write an R data set, we'll use `write_rds(datasetname, "locationoncomputer")`. The function `saveRDS()` would also work, in a similar way, but be a little slower.

# The dm1000 data

```
dm1000
```

```
# A tibble: 1,000 x 17
```

	subject	sbp	dbp	insurance	age	n_income	ht
	<chr>	<dbl>	<dbl>	<fct>	<dbl>	<dbl>	<dbl>
1	M-0001	145	70	Medicaid	55	29853	1.63
2	M-0002	151	77	Commercial	52	31248	1.75
3	M-0003	127	73	Medicare	69	23362	1.65
4	M-0004	125	74	Medicaid	57	26033	1.63
5	M-0005	120	73	Medicare	68	85374	1.69
6	M-0006	127	75	Medicaid	56	31273	1.71
7	M-0007	114	81	Commercial	54	25445	1.68
8	M-0008	166	110	Medicare	45	67526	1.69
9	M-0009	111	77	Medicare	61	15203	1.91
10	M-0010	146	102	Medicaid	63	17628	1.86

```
# ... with 990 more rows, and 10 more variables:
```

```
#   wt <dbl>, a1c <dbl>, ldl <dbl>, tobacco <fct>,
```

```
#   statin <dbl>, eye_exam <dbl>
```

# Partitioning dm1000 into two groups

Before we do anything else today, let's split the data in dm1000 who have complete data on sbp and dbp into two groups:

- a model **development** or **training** sample (70% of rows)
- a model **evaluation** or **test** sample (the other 30%)

There are many ways to do this in R. Let's start by filtering out the observations with missing values of blood pressure.

```
dm994 <- dm1000 %>% filter(complete.cases(sbp, dbp)) %>%  
  select(subject, sbp, dbp, insurance)
```

```
dm994 %>% nrow()
```

```
[1] 994
```

```
dm994 %$% n_distinct(subject)
```

```
[1] 994
```

# Now, let's build the partition.

Again, we want 70% of the sample in our training set, and the remaining 30% in our test set.

```
set.seed(4312021) # for replicating the sampling later

dm_train <- dm994 %>% sample_frac(0.7)
dm_test  <- dm994 %>% anti_join(dm_train, by = "subject")

nrow(dm_train); nrow(dm_test)
```

```
[1] 696
```

```
[1] 298
```

OK. Looks good!




Can dbp predict sbp?

# Plotting sbp vs. dbp (training set)

Positive Association of SBP and DBP

loess smooth in blue, OLS model in red



696 subjects from dm\_train.

- Note caption, labels, increased text size for Pearson r.


## Code from Previous Slide

```
ggplot(data = dm_train, aes(x = dbp, y = sbp)) +  
  geom_point() +  
  geom_smooth(method = "loess", col = "blue",  
              se = FALSE, formula = y ~ x) +  
  geom_smooth(method = "lm", col = "red",  
              se = FALSE, formula = y ~ x) +  
  geom_label(x = 120, y = 100, size = 5,  
            label = glue('Pearson r = {round_half_up(  
  cor(dm_train$sbp, dm_train$dbp),2)}.')) +  
  geom_label_repel(data = dm_train %>% filter(dbp > 120),  
                  aes(label = subject), fill = "yellow") +  
  labs(title = "Positive Association of SBP and DBP",  
       subtitle = "loess smooth in blue, OLS model in red",  
       caption =  
         glue('{nrow(dm_train)} subjects from dm_train.'))
```

# Redo plot without point M-0862

What happens if we drop M-0862?

Newly fit loess in blue, new OLS in red



695 subjects.

- Note increased size and new color of point M-0862, use of `geom_text_repel` instead of `geom_label_repel`, adjusted caption.

## Code from Previous Slide

```
ggplot(data = dm_train, aes(x = dbp, y = sbp)) +  
  geom_point() +  
  geom_smooth(data = dm_train %>% filter(dbp <= 120),  
              method = "loess", col = "blue",  
              se = FALSE, formula = y ~ x) +  
  geom_smooth(data = dm_train %>% filter(dbp <= 120),  
              method = "lm", col = "red",  
              se = FALSE, formula = y ~ x) +  
  geom_point(data = dm_train %>% filter(dbp > 120),  
             aes(col = "purple", size = 3)) +  
  geom_text_repel(data = dm_train %>% filter(dbp > 120),  
                  aes(label = subject)) +  
  guides(color = "none", size = "none") +  
  labs(title = "What happens if we drop M-0862?",  
        subtitle = "Newly fit loess in blue, new OLS in red",  
        caption = glue('{nrow(dm_train)-1} subjects.'))
```

# Modeling sbp using dbp (training set)

```
m1_train <- lm(sbp ~ dbp, data = dm_train)
```

```
tidy(m1_train, conf.int = TRUE, conf.level = 0.90) %>%  
  select(term, estimate, conf.low, conf.high) %>% kable()
```

term	estimate	conf.low	conf.high
(Intercept)	80.6798905	74.4662421	86.893539
dbp	0.6982168	0.6160396	0.780394

```
glance(m1_train) %>% select(nobs, r.squared) %>% kable()
```

nobs	r.squared
696	0.2200811

# Summarizing the Training Fit

- 1 We can use `extract_eq()` from the `equatiomatic` package to present the equation from our model in a fairly attractive way, but we must use the code chunk header `{r, results = 'asis'}`.

```
extract_eq(m1_train, use_coefs = TRUE, coef_digits = 3)
```

$$\widehat{\text{sbp}} = 80.68 + 0.698(\text{dbp})$$

- 2 The `summary` function when applied to a linear model (`lm`) produces output that isn't organized in a way that allows up to plot or present it effectively outside of an R session.

```
summary(m1_train)
```

A screenshot follows on the next page.

```

> summary(m1_train)

Call:
lm(formula = sbp ~ dbp, data = dm_train)

Residuals:
    Min       1Q   Median       3Q      Max
-37.159 -11.537  -1.348   10.066   52.990

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  80.67989    3.77259   21.39  <2e-16 ***
dbp           0.69822    0.04989   13.99  <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 16.12 on 694 degrees of freedom
Multiple R-squared:  0.2201,    Adjusted R-squared:  0.219
F-statistic: 195.8 on 1 and 694 DF,  p-value: < 2.2e-16

```



# Why I like tidy() and other broom functions

**broom:** turn messy model outputs into tidy TIBBLES!



@allison\_horst

<https://github.com/allisonhorst/stats-illustrations>

# Does R like this linear model?

```
tidy(m1_train) %>% kable(digits = 3)
```

term	estimate	std.error	statistic	p.value
(Intercept)	80.680	3.773	21.386	0
dbp	0.698	0.050	13.994	0

Yes. Wow. It **really** does. Look at those  $p$  values!

# How much of the variation in sbp does m1 capture?

The glance function can help us (again from broom.)

```
glance(m1_train) %>%  
  select(nobs, r.squared, p.value, sigma) %>% kable()
```

nobs	r.squared	p.value	sigma
696	0.2200811	0	16.11605


- $r.squared = R^2$ , the proportion of variation in sbp accounted for by the model using dbp.
  - indicates improvement over predicting  $\text{mean}(sbp)$  for everyone
- $p.value$  = refers to a global F test
  - indicates something about combination of  $r^2$  and sample size
- $\sigma$  = residual standard error

glance provides 9 additional summaries for a linear model.


# Calibrating Yourself on R-square

# Can you match each plot to its R-square?


A. R-square = ?




B. R-square = ?




C. R-square = ?




D. R-square = ?



E. R-square = ?









F. R-sq = ?



- $R^2$  values shown include 0.16, 0.25, 0.49, 0.64, 0.81 and 0.91

# Gaining Insight into what R-square implies



# Obtaining Residuals and Fitted Values in the Training Sample

## Predict using m1\_train: $\text{sbp} = 80.68 + 0.70 \text{ dbp}$

Use `augment` (from `broom`) to capture fitted values and residuals for all of the data in the training sample.

```
augment(m1_train, data = dm_train) %>%  
  select(subject, sbp, dbp, .fitted, .resid) %>%  
  slice_min(., order_by = subject, n = 2) %>% kable(dig = 2)
```

subject	sbp	dbp	.fitted	.resid
M-0002	151	77	134.44	16.56
M-0003	127	73	131.65	-4.65

- Subject M-0002 has an observed sbp of 151, and dbp of 77.
- Our `m1_train` model **fits** (predicts) M-0002's sbp to be 134.44, so that's a **residual** of  $151 - 134.44 = 16.56$  mm Hg.
- Note that **residual** = **observed** - **fitted**.




# What must we assume for a regression model?

Briefly (for now), we assume that:


- the regression relationship is linear, rather than curved, and we can assess this by plotting the regression residuals (prediction errors) against the fitted values and looking to see if a curve emerges
- the regression residuals show similar variance across levels of the fitted values, and again we can get insight into this by plotting residuals vs. predicted values
- the regression residuals (prediction errors) are well described by a Normal model, and we can assess this with all of our usual visualizations to help decide on whether a Normal model is reasonable for a batch of data.
- We assess all of these issues (and others) with plots of the residuals. Let's start with the **Normality** assumption. . .

# Plot residuals from m1\_train


Normal Q-Q: 994 m1 Residuals



Hist + Normal Density: m1 Residuals



Boxplot: m1 Residuals



min	Q1	median	Q3	max	mean	sd	n	missing
-37.2	-11.5	-1.3	10.1	53	0	16.1	696	0

# Plot Residuals vs. Predicted (Fitted) Values (code)


```
m1_train_aug <- augment(m1_train, data = dm_train)

ggplot(m1_train_aug, aes(x = .fitted, y = .resid)) +
  geom_point() +
  geom_smooth(method = "lm", col = "red",
              formula = y ~ x, se = FALSE) +
  geom_smooth(method = "loess", col = "blue",
              formula = y ~ x, se = FALSE) +
  labs(title = "m1_train: Residuals vs. Fitted Values",
       x = "Fitted sbp values", y = "Residuals")
```

# m1\_train: Residuals vs. Predicted (Fitted) Values

- We're looking to see if there is a substantial curve in the plot, or if the variability changes materially from left to right.
- What we want to see is a “fuzzy football” actually.

m1\_train: Residuals vs. Fitted Values



This sort of fuzzy football...



# Making Predictions Out of Sample (into the Test Sample)

# Use model `m1_train` to predict SBP in `dm_test`

```
m1_test_aug <- augment(m1_train, newdata = dm_test)
```

```
m1_test_aug %>% nrow()
```

```
[1] 298
```


- We have predictions from `m1_train` for the 298 subjects in `dm_test`.
- Remember we didn't use the `dm_test` data to build `m1_train`.

```
m1_test_aug %>%  
  select(subject, sbp, dbp, .fitted, .resid) %>%  
  slice_min(., order_by = subject, n = 2) %>% kable(dig = 2)
```


subject	sbp	dbp	.fitted	.resid
M-0001	145	70	129.56	15.44
M-0007	114	81	137.24	-23.24

# dm\_test (n = 298): m1\_train Prediction Errors


Normal Q-Q: 298 m1\_train Errors



Hist + Normal Density: 298 m1\_train Errors



Boxplot: 298 m1\_train Errors



min	Q1	median	Q3	max	mean	sd	n	missing
-43.1	-10.2	-1	9	71.1	0.3	15.9	298	0



# Out-of-Sample (Test Set) Error Summaries (m1)

- Mean Absolute Prediction Error = 12.14
- Maximum Absolute Prediction Error = 71.07
- (square Root of) Mean Squared Prediction Error (RMSPE) = 15.83

```
mosaic::favstats(~ abs(.resid), data = m1_test_aug) %>%  
  select(n, min, median, max, mean, sd) %>%  
  kable(digits = 2)
```

n	min	median	max	mean	sd
298	0.03	9.82	71.07	12.14	10.18


```
sqrt(mean(m1_test_aug$.resid^2))
```

```
[1] 15.83017
```

These statistics are most useful when we're comparing two models.

## Back to all 994 values. Does m1\_train work well?

```
ggplot(dm994, aes(x = dbp, y = sbp)) +  
  geom_point() + theme(aspect.ratio = 1) +  
  geom_abline(intercept = 80.6799, slope = 0.6982,  
             col = "purple", lwd = 2.5, alpha = 0.5)
```



# Is this the only linear model R can fit to these data?

Nope.

```
library(rstanarm)
```

```
m2_train <- stan_glm(sbp ~ dbp, data = dm_train)
```

```
SAMPLING FOR MODEL 'continuous' NOW (CHAIN 1).
```

```
Chain 1:
```

```
Chain 1: Gradient evaluation took 0 seconds
```

```
Chain 1: 1000 transitions using 10 leapfrog steps per transition
```

```
Chain 1: Adjust your expectations accordingly!
```

```
Chain 1:
```

```
Chain 1:
```

```
Chain 1: Iteration:      1 / 2000 [  0%] (Warmup)
```

```
Chain 1: Iteration:    200 / 2000 [ 10%] (Warmup)
```

```
Chain 1: Iteration:    400 / 2000 [ 20%] (Warmup)
```

```
Chain 1: Iteration:    600 / 2000 [ 30%] (Warmup)
```

```
Chain 1: Iteration:    800 / 2000 [ 40%] (Warmup)
```

# Bayesian fitted linear model for our sbp data

```
print(m2_train)
```

```
stan_glm
  family:      gaussian [identity]
 formula:      sbp ~ dbp
 observations: 696
 predictors:    2
```

-----

	Median	MAD_SD
(Intercept)	80.8	3.7
dbp	0.7	0.0

Auxiliary parameter(s):

	Median	MAD_SD
sigma	16.1	0.4

-----

# Is the Bayesian model (with default prior) very different from our `lm` in this situation?

```
broom::tidy(m1_train) # fit with lm
```

```
# A tibble: 2 x 5
```

	term	estimate	std.error	statistic	p.value
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	(Intercept)	80.7	3.77	21.4	2.47e-78
2	dbp	0.698	0.0499	14.0	2.23e-39

```
broom.mixed::tidy(m2_train) # stan_glm with default priors
```

```
# A tibble: 2 x 3
```

	term	estimate	std.error
	<chr>	<dbl>	<dbl>
1	(Intercept)	80.8	3.74
2	dbp	0.697	0.0490

# Test Sample fits and residuals from Bayesian model

```
m2_test_aug <- dm_test %>% select(subject, sbp, dbp) %>%  
  mutate(.fitted = predict(m2_train, newdata = dm_test),  
         .resid = sbp - .fitted)  
  
m2_test_aug %>%  
  select(subject, sbp, dbp, .fitted, .resid) %>%  
  slice_min(., order_by = subject, n = 2) %>% kable(dig = 2)
```

subject	sbp	dbp	.fitted	.resid
M-0001	145	70	129.56	15.44
M-0007	114	81	137.23	-23.23

# Out-of-Sample (Test Set) Error Summaries (m2)

```
mosaic::favstats(~ abs(.resid), data = m2_test_aug) %>%  
  select(n, min, median, max, mean, sd) %>%  
  kable(digits = 3)
```

n	min	median	max	mean	sd
298	0.021	9.807	71.071	12.137	10.178

```
sqrt(mean(m2_test_aug$.resid^2))
```

```
[1] 15.82868
```


Test Set Error Summary	OLS model m1	Bayes model m2
Mean Absolute Prediction Error	12.139	12.137
Maximum Absolute Prediction Error	71.066	71.071
Root Mean Squared Prediction Error	15.83	15.829

What if we add another predictor? (Insurance)



# Plotting sbp vs. dbp and insurance

```
ggplot(data = dm_train, aes(x = dbp, y = sbp,  
                             col = insurance, group = insurance)) +  
  geom_point() +  
  geom_smooth(method = "lm", formula = y ~ x, se = FALSE)
```



# Two possible models

```
m3_train <- lm(sbp ~ dbp + insurance, data = dm_train)
m4_train <- lm(sbp ~ dbp * insurance, data = dm_train)
```

- What is the difference between m3 and m4?
  - Model m3 will allow the intercept term of the sbp-dbp relationship to vary depending on insurance.
  - Model m4 will allow both the slope and intercept of the sbp-dbp relationship to vary depending on insurance.

## Equation for m3 (sbp ~ dbp + insurance)

```
extract_eq(m3_train, use_coefs = TRUE,  
           wrap = TRUE, terms_per_line = 2)
```

$$\widehat{\text{sbp}} = 77.58 + 0.72(\text{dbp}) + \\ 1.11(\text{insurance}_{\text{Commercial}}) + 2.73(\text{insurance}_{\text{Medicare}}) + \\ 1.16(\text{insurance}_{\text{Uninsured}})$$

- Predicted sbp by m3 for a Commercial subject?

## Equation for m3 (sbp ~ dbp + insurance)

```
extract_eq(m3_train, use_coefs = TRUE,  
           wrap = TRUE, terms_per_line = 2)
```

$$\widehat{\text{sbp}} = 77.58 + 0.72(\text{dbp}) + \\ 1.11(\text{insurance}_{\text{Commercial}}) + 2.73(\text{insurance}_{\text{Medicare}}) + \\ 1.16(\text{insurance}_{\text{Uninsured}})$$

- Predicted sbp by m3 for a Commercial subject?
- $\text{sbp} = 77.58 + 0.72 * \text{dbp} + 1.11(1) + 2.73(0) + 1.16(0)$

## Equation for m3 (sbp ~ dbp + insurance)

```
extract_eq(m3_train, use_coefs = TRUE,  
           wrap = TRUE, terms_per_line = 2)
```

$$\widehat{\text{sbp}} = 77.58 + 0.72(\text{dbp}) + \\ 1.11(\text{insurance}_{\text{Commercial}}) + 2.73(\text{insurance}_{\text{Medicare}}) + \\ 1.16(\text{insurance}_{\text{Uninsured}})$$

- Predicted sbp by m3 for a Commercial subject?
- $\text{sbp} = 77.58 + 0.72 * \text{dbp} + 1.11(1) + 2.73(0) + 1.16(0)$
- $\text{sbp} = 78.69 + 0.72 * \text{dbp}$

## Equation for m3 (sbp ~ dbp + insurance)

```
extract_eq(m3_train, use_coefs = TRUE,  
           wrap = TRUE, terms_per_line = 2)
```

$$\widehat{\text{sbp}} = 77.58 + 0.72(\text{dbp}) + \\ 1.11(\text{insurance}_{\text{Commercial}}) + 2.73(\text{insurance}_{\text{Medicare}}) + \\ 1.16(\text{insurance}_{\text{Uninsured}})$$

- Predicted sbp by m3 for a Commercial subject?
- $\text{sbp} = 77.58 + 0.72 * \text{dbp} + 1.11(1) + 2.73(0) + 1.16(0)$
- $\text{sbp} = 78.69 + 0.72 * \text{dbp}$
- For a Medicaid subject, m3 predicts  $\text{sbp} = 77.58 + 0.72 \text{ dbp}$

## Equation for m3 (sbp ~ dbp + insurance)

```
extract_eq(m3_train, use_coefs = TRUE,  
           wrap = TRUE, terms_per_line = 2)
```

$$\widehat{\text{sbp}} = 77.58 + 0.72(\text{dbp}) + \\ 1.11(\text{insurance}_{\text{Commercial}}) + 2.73(\text{insurance}_{\text{Medicare}}) + \\ 1.16(\text{insurance}_{\text{Uninsured}})$$

- Predicted sbp by m3 for a Commercial subject?
- $\text{sbp} = 77.58 + 0.72 \cdot \text{dbp} + 1.11(1) + 2.73(0) + 1.16(0)$
- $\text{sbp} = 78.69 + 0.72 \cdot \text{dbp}$
- For a Medicaid subject, m3 predicts  $\text{sbp} = 77.58 + 0.72 \cdot \text{dbp}$
- For a Medicare subject, m3 predicts  $\text{sbp} = 80.31 + 0.72 \cdot \text{dbp}$

## Equation for m3 (sbp ~ dbp + insurance)

```
extract_eq(m3_train, use_coefs = TRUE,  
           wrap = TRUE, terms_per_line = 2)
```

$$\widehat{\text{sbp}} = 77.58 + 0.72(\text{dbp}) + \\ 1.11(\text{insurance}_{\text{Commercial}}) + 2.73(\text{insurance}_{\text{Medicare}}) + \\ 1.16(\text{insurance}_{\text{Uninsured}})$$

- Predicted sbp by m3 for a Commercial subject?
- $\text{sbp} = 77.58 + 0.72 * \text{dbp} + 1.11(1) + 2.73(0) + 1.16(0)$
- $\text{sbp} = 78.69 + 0.72 * \text{dbp}$
- For a Medicaid subject, m3 predicts  $\text{sbp} = 77.58 + 0.72 \text{ dbp}$
- For a Medicare subject, m3 predicts  $\text{sbp} = 80.31 + 0.72 \text{ dbp}$
- For an uninsured subject, m3 predicts  $\text{sbp} = 78.74 + 0.72 \text{ dbp}$



## Equation for m3 (sbp ~ dbp + insurance)

```
extract_eq(m3_train, use_coefs = TRUE,  
           wrap = TRUE, terms_per_line = 2)
```

$$\widehat{\text{sbp}} = 77.58 + 0.72(\text{dbp}) + \\ 1.11(\text{insurance}_{\text{Commercial}}) + 2.73(\text{insurance}_{\text{Medicare}}) + \\ 1.16(\text{insurance}_{\text{Uninsured}})$$

- Predicted sbp by m3 for a Commercial subject?
- $\text{sbp} = 77.58 + 0.72 \cdot \text{dbp} + 1.11(1) + 2.73(0) + 1.16(0)$
- $\text{sbp} = 78.69 + 0.72 \cdot \text{dbp}$
- For a Medicaid subject, m3 predicts  $\text{sbp} = 77.58 + 0.72 \cdot \text{dbp}$
- For a Medicare subject, m3 predicts  $\text{sbp} = 80.31 + 0.72 \cdot \text{dbp}$
- For an uninsured subject, m3 predicts  $\text{sbp} = 78.74 + 0.72 \cdot \text{dbp}$
- Note: only the intercept term varies by insurance in m3.

## Equation for m4 (sbp ~ dbp \* insurance)

```
extract_eq(m4_train, use_coefs = TRUE,  
           wrap = TRUE, terms_per_line = 2)
```

$$\begin{aligned}\widehat{\text{sbp}} = & 60.26 + 0.94(\text{dbp}) + \\ & 23.54(\text{insurance}_{\text{Commercial}}) + 31.04(\text{insurance}_{\text{Medicare}}) + \\ & 25.78(\text{insurance}_{\text{Uninsured}}) - 0.29(\text{dbp} \times \text{insurance}_{\text{Commercial}}) - \\ & 0.38(\text{dbp} \times \text{insurance}_{\text{Medicare}}) - 0.32(\text{dbp} \times \text{insurance}_{\text{Uninsured}})\end{aligned}$$

- m4 predicts, for a Commercial subject. . .

## Equation for m4 (sbp ~ dbp \* insurance)

```
extract_eq(m4_train, use_coefs = TRUE,  
           wrap = TRUE, terms_per_line = 2)
```

$$\begin{aligned}\widehat{\text{sbp}} = & 60.26 + 0.94(\text{dbp}) + \\ & 23.54(\text{insurance}_{\text{Commercial}}) + 31.04(\text{insurance}_{\text{Medicare}}) + \\ & 25.78(\text{insurance}_{\text{Uninsured}}) - 0.29(\text{dbp} \times \text{insurance}_{\text{Commercial}}) - \\ & 0.38(\text{dbp} \times \text{insurance}_{\text{Medicare}}) - 0.32(\text{dbp} \times \text{insurance}_{\text{Uninsured}})\end{aligned}$$

- m4 predicts, for a Commercial subject. . .
- $\text{sbp} = 60.26 + 0.94 * \text{dbp} + 23.54 (1) + 31.04 (0) + 25.78 (0) - 0.29 (\text{dbp} * 1) - 0.38 (\text{dbp} * 0) - 0.32 (\text{dbp} * 0)$

## Equation for m4 (sbp ~ dbp \* insurance)

```
extract_eq(m4_train, use_coefs = TRUE,  
           wrap = TRUE, terms_per_line = 2)
```

$$\begin{aligned}\widehat{\text{sbp}} = & 60.26 + 0.94(\text{dbp}) + \\ & 23.54(\text{insurance}_{\text{Commercial}}) + 31.04(\text{insurance}_{\text{Medicare}}) + \\ & 25.78(\text{insurance}_{\text{Uninsured}}) - 0.29(\text{dbp} \times \text{insurance}_{\text{Commercial}}) - \\ & 0.38(\text{dbp} \times \text{insurance}_{\text{Medicare}}) - 0.32(\text{dbp} \times \text{insurance}_{\text{Uninsured}})\end{aligned}$$

- m4 predicts, for a Commercial subject. . .
- $\text{sbp} = 60.26 + 0.94 * \text{dbp} + 23.54 (1) + 31.04 (0) + 25.78 (0) - 0.29 (\text{dbp} * 1) - 0.38 (\text{dbp} * 0) - 0.32 (\text{dbp} * 0)$
- $\text{sbp} = (60.26 + 23.54) + (0.94 - 0.29) * \text{dbp}$

## Equation for m4 (sbp ~ dbp \* insurance)

```
extract_eq(m4_train, use_coefs = TRUE,  
           wrap = TRUE, terms_per_line = 2)
```

$$\begin{aligned}\widehat{\text{sbp}} = & 60.26 + 0.94(\text{dbp}) + \\ & 23.54(\text{insurance}_{\text{Commercial}}) + 31.04(\text{insurance}_{\text{Medicare}}) + \\ & 25.78(\text{insurance}_{\text{Uninsured}}) - 0.29(\text{dbp} \times \text{insurance}_{\text{Commercial}}) - \\ & 0.38(\text{dbp} \times \text{insurance}_{\text{Medicare}}) - 0.32(\text{dbp} \times \text{insurance}_{\text{Uninsured}})\end{aligned}$$

- m4 predicts, for a Commercial subject. . .
- $\text{sbp} = 60.26 + 0.94 * \text{dbp} + 23.54 (1) + 31.04 (0) + 25.78 (0) - 0.29 (\text{dbp} * 1) - 0.38 (\text{dbp} * 0) - 0.32 (\text{dbp} * 0)$
- $\text{sbp} = (60.26 + 23.54) + (0.94 - 0.29) * \text{dbp}$
- $\text{sbp} = 83.80 - 0.65 \text{ dbp}$  for Commercial subjects

## Equation for m4 (sbp ~ dbp \* insurance)

```
extract_eq(m4_train, use_coefs = TRUE,  
           wrap = TRUE, terms_per_line = 2)
```

$$\widehat{\text{sbp}} = 60.26 + 0.94(\text{dbp}) + 23.54(\text{insurance}_{\text{Commercial}}) + 31.04(\text{insurance}_{\text{Medicare}}) + 25.78(\text{insurance}_{\text{Uninsured}}) - 0.29(\text{dbp} \times \text{insurance}_{\text{Commercial}}) - 0.38(\text{dbp} \times \text{insurance}_{\text{Medicare}}) - 0.32(\text{dbp} \times \text{insurance}_{\text{Uninsured}})$$

- For Medicaid subjects,  $\text{sbp} = 60.26 + 0.94 * \text{dbp}$
- For Medicare subjects,  $\text{sbp} = 91.30 + 0.56 * \text{dbp}$
- For the uninsured,  $\text{sbp} = 86.04 + 0.62 * \text{dbp}$
- So both the slope and the intercept are changing in m4

# How do these models do in the training sample?

- Model m3

```
glance(m3_train) %>%  
  select(r.squared, adj.r.squared, sigma, AIC, BIC) %>%  
  kable(digits = c(3, 3, 1, 1, 1))
```

r.squared	adj.r.squared	sigma	AIC	BIC
0.224	0.22	16.1	5851.1	5878.4

- Model m4

```
glance(m4_train) %>%  
  select(r.squared, adj.r.squared, sigma, AIC, BIC) %>%  
  kable(digits = c(3, 3, 1, 1, 1))
```

r.squared	adj.r.squared	sigma	AIC	BIC
0.236	0.229	16	5846	5886.9

# Augmenting and Testing Models m3 and m4


```
m3_train_aug <- augment(m3_train, data = dm_train)
m3_test_aug <- augment(m3_train, newdata = dm_test)

m4_train_aug <- augment(m4_train, data = dm_train)
m4_test_aug <- augment(m4_train, newdata = dm_test)
```




# Residuals (training sample) for m3\_train


Normal Q-Q: 698 m3 Residuals



Hist + Normal Density: m3 Residuals



Boxplot: m3 Residuals




min	Q1	median	Q3	max	mean	sd	n	missing
-38.3	-11.5	-1.3	9.8	54	0	16.1	696	0

# m3\_train: Residuals vs. Predicted (Fitted) Values


- We're looking for a “fuzzy football”...

m3\_train: Residuals vs. Fitted Values




# Residuals (training sample) for m4\_train

Normal Q-Q: 698 m4 Residuals



Hist + Normal Density: m4 Residuals



Boxplot: m4 Residuals




min	Q1	median	Q3	max	mean	sd	n	missing
-41.2	-11.1	-1.3	9.6	50.1	0	15.9	696	0

# m4\_train: Residuals vs. Predicted (Fitted) Values

- We're looking for a “fuzzy football”...

m4\_train: Residuals vs. Fitted Values



# Comparing performance on the training data

```
bind_rows(glance(m1_train), glance(m2_train),  
          glance(m3_train), glance(m4_train)) %>%  
mutate(modname = c("m1", "m2", "m3", "m4")) %>%  
select(modname, r2 = r.squared, adj_r2 = adj.r.squared,  
       sigma, AIC, BIC) %>%  
kable(digits = c(0, 3, 3, 2, 1, 1))
```

modname	r2	adj_r2	sigma	AIC	BIC
m1	0.220	0.219	16.12	5848.7	5862.3
m2	NA	NA	16.12	NA	NA
m3	0.224	0.220	16.11	5851.1	5878.4
m4	0.236	0.229	16.02	5846.0	5886.9

- The `glance()` function produces different results for a Bayesian `stan_glm()` model like `m2`, so we'll ignore that for now.

# Comparing performance on the test data

Here are some fundamental summaries of absolute prediction error (APE) along with the root mean squared prediction error (RMSPE) for each of our models, in the **testing** sample.

Summary	Mean APE	Max APE	RMSPE
m1_train: lm	12.139	71.066	15.83
m2_train: stan_glm	12.137	71.071	15.829
m3_train: dbp+insurance	12.04	72.367	15.778
m4_train: dbp*insurance	11.947	71.37	15.647

- Which of these models displays the strongest predictive performance in our test sample?

# Reminder of Today's Agenda

- 1 Ingesting `dm1000` data using R data set format (`.Rds`)
- 2 Partitioning data into model training/test samples.
- 3 Augmenting a Scatterplot (labeling, size, color) and fitting a simple OLS (linear) model `m1`
- 4 Using `summary()` and `extract_eq()` on a regression model.
- 5 The broom package and `tidy()`, `glance()` and `augment()`
- 6 Calibrating your understanding of R-square a bit
- 7 Assessing Regression Assumptions with Residual Plots
- 8 Making Predictions into the Test Sample
- 9 Assessing Quality of Fit using the Test Sample with mean and maximum absolute prediction error and with RMSPE
- 10 Fitting a Bayesian Linear Model with default priors (`m2`)
- 11 Including Insurance without (`m3`) and with (`m4`) interaction with `dbp` in linear models